


# Introduction to Computing and Programming in Python™: A Multimedia Approach

Fourth Edition



## Chapter 4

### Modifying Pictures Using Loops

# Learning Objectives (1 of 2)

- 4.1** To understand how images are digitized by taking advantage of limits in human vision.
- 4.2** To identify different models for color, including RGB, the most common one for computers.
- 4.3** To manipulate color values in pictures, like increasing or decreasing red values.
- 4.4** To convert a color picture to grayscale, using more than one method.
- 4.5** To negate a picture.

# Learning Objectives (2 of 2)

**4.6** To use a matrix representation in finding pixels in a picture.

**4.7** To use the objects **pictures** and **pixels**.

**4.8** To use iteration (with a for loop) for changing the color values of pixels in a picture.

**4.9** To nest blocks of code within one another.

**4.10** To choose between having a function **return** a value and just providing a **side effect**.

**4.11** To determine the **scope** of a variable name.

# We Perceive Light Different from How it Actually is

- Color is continuous
  - Visible light is in the wavelengths between 370 and 730 nanometers
    - That's 0.00000037 and 0.00000073 meters
- But we **perceive** light with color sensors that peak around 425 nm (blue), 550 nm (green), and 560 nm (red).
  - Our brain figures out which color is which by figuring out how much of each kind of sensor is responding
  - One implication: We perceive two kinds of “orange” — one that's **spectral** and one that's red + yellow (hits our color sensors just right)
  - Dogs and other simpler animals have only two kinds of sensors
    - They **do** see color. Just **less** color.

# Luminance Vs. Color


- We perceive **borders** of things, motion, depth via **luminance**
  - Luminance is **not** the amount of light, but our **perception** of the amount of light.
  - We see blue as “darker” than red, even if same amount of light.
- Much of our luminance perception is based on comparison to backgrounds, not raw values.



Luminance is actually **color blind**. Completely different part of the brain does luminance vs. color.

# Digitizing Pictures as Bunches of Little Dots

- We digitize pictures into lots of little dots
- Enough dots and it looks like a continuous whole to our eye
  - Our eye has limited resolution
  - Our background/depth **acuity** is particularly low
- Each picture element is referred to as a **pixel**

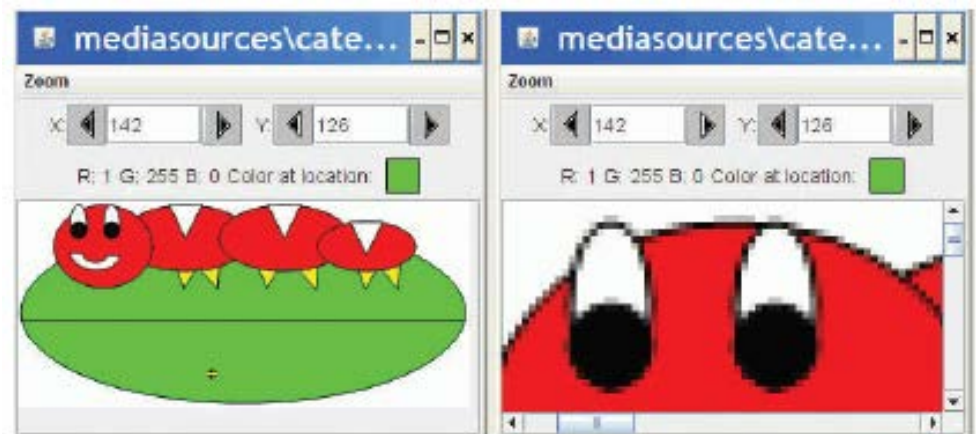


# Pixels

- Pixels are **picture elements**
  - Each pixel object knows its **color**
  - It also knows where it is in its **picture**

```
>>> file = "c:/ip-book/mediasources/caterpillar.jpg"  
>>> pict = makePicture(file)  
>>> explore(pict)
```

When we zoom the picture to 500%, we can see individual pixels.



# A Picture is a Matrix of Pixels

- It's not a continuous line of elements, that is, an **array**
- A picture has two dimensions: Width and Height
- We need a two-dimensional array: a **matrix**

	0	1	2	3
	15	12	13	10

	0	1	2	3
0	15	12	13	10
1	9	7	2	1
2	6	3	9	10



# Referencing a Matrix

- We talk about positions in a matrix as (x,y), or (horizontal, vertical)
- Element (1,0) in the matrix at left is the value 12
- Element (0,2) is 6





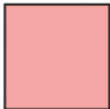
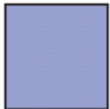


	0	1	2	3
0	15	12	13	10
1	9	7	2	1
2	6	3	9	10

# Encoding Color

- Each pixel encodes color at that position in the picture
- Lots of encodings for color
  - Printers use CMYK: Cyan, Magenta, Yellow, and black.
  - Others use HSB for Hue, Saturation, and Brightness (also called HSV for Hue, Saturation, and Value)
- We'll use the most common for computers
  - RGB: Red, Green, Blue

# Encoding RGB

- Each component color (red, green, and blue) is encoded as a single byte
- Colors go from (0,0,0) to (255,255,255)
  - If all three components are the same, the color is in greyscale
    - (200,200,200) at (3,1)
  - (0,0,0) (at position (3,0) in example) is black
  - (255,255,255) is white

	0	1	2	3
0	 255, 30, 30	 30, 30, 255	 30, 255, 30	 0, 0, 0
1	 255, 150, 150	 150, 150, 255	 150, 255, 150	 200, 200, 200

# How Much Can We Encode in 8 Bits?

- Let's walk it through.
  - If we have one bit, we can represent two patterns: 0 and 1.
  - If we have two bits, we can represent four patterns:  
00, 01, 10, 11.
  - If we have three bits, we can represent eight patterns:  
000, 001, 010, 011, 100, 101, 110, 111
- General rule: In  $n$  bits, we can have  $2^n$  patterns
  - In 8 bits, we can have  $2^8$  patterns, or 256
  - If we make one pattern 0, then the highest value we can represent is  $2^8 - 1$  or 255

# Is That Enough?

- We're representing color in  $24(3 * 8)$  bits.
  - That's 16,777,216 ( $2^{24}$ ) possible colors
  - Our eye can discern millions of colors, so it's probably pretty close
  - But the real limitation is the physical devices: We don't get 16 million colors out of a monitor
- Some graphics systems support 32 bits per pixel
  - May be more pixels for color, or an additional 8 bits to represent 256 levels of **translucence**



# Size of Images

	<b>320 x 240 image</b>	<b>640 x 480 image</b>	<b>1024 x 768 monitor</b>
<b>24 bit color</b>	230,400 bytes	921,600 bytes	2,359,296 bytes
<b>32 bit color</b>	307,200 bytes	1,228,800 bytes	3,145,728 bytes